

hw09

May 19, 2022

Name: Allan Gongora

Section: 1031

1 Homework 9: Central Limit Theorem

Please complete this notebook by filling in the cells provided. Before you begin, execute the following cell to load the provided tests.

```
[1]: pip install gofer-grader
```

```
Requirement already satisfied: gofer-grader in /opt/conda/lib/python3.7/site-  
packages (1.1.0)  
Requirement already satisfied: tornado in /opt/conda/lib/python3.7/site-packages  
(from gofer-grader) (6.1)  
Requirement already satisfied: jinja2 in /opt/conda/lib/python3.7/site-packages  
(from gofer-grader) (3.0.3)  
Requirement already satisfied: pygments in /opt/conda/lib/python3.7/site-  
packages (from gofer-grader) (2.11.2)  
Requirement already satisfied: MarkupSafe>=2.0 in /opt/conda/lib/python3.7/site-  
packages (from jinja2->gofer-grader) (2.0.1)  
Note: you may need to restart the kernel to use updated packages.
```

```
[2]: # Don't change this cell; just run it.  
  
import numpy as np  
from datascience import *  
  
# These lines do some fancy plotting magic.  
import matplotlib  
%matplotlib inline  
import matplotlib.pyplot as plt  
plt.style.use('fivethirtyeight')  
import warnings  
warnings.simplefilter('ignore', FutureWarning)  
  
# These lines load the tests.
```

```
from gofer.ok import check
```

Recommended Reading: * [Why the Mean Matters](#)

- 1) For all problems that you must write explanations and sentences for, you **must** provide your answer in the designated space. This can include:
 - A) Sentence responses to questions that ask for an explanation
 - B) Numeric responses to multiple choice questions
 - C) Programming code
- 2) Moreover, throughout this homework and all future ones, please be sure to not re-assign variables throughout the notebook! For example, if you use `max_temperature` in your answer to one question, do not reassign it later on. Otherwise, you will fail tests that you thought you were passing previously!

Once you're finished, select "Save and Checkpoint" in the File menu. Your name and course section number should be in the first and last cell of the assignment. Be sure you have run all cells with code and that the output from that is showing. Then click "Print Preview" in the File menu. Print a copy from there in pdf format. (This means you right click and choose print and choose "save as pdf" from your printer options.) You will need to submit the pdf in Canvas by the deadline.

The gopher grader output and/or output from your coding are essential to helping your instructor grade your work correctly and in a timely manner.

Files submitted that are missing the required output will lose some to all points so double check your pdf before submitting.

1.1 1. The Bootstrap and The Normal Curve

In this exercise, we will explore a dataset that includes the safety inspection scores for restaurants in the city of Austin, Texas. We will be interested in determining the average restaurant score for the city from a random sample of the scores; the average restaurant score is out of 100. We'll compare two methods for computing a confidence interval for that quantity: the bootstrap resampling method, and an approximation based on the Central Limit Theorem.

```
[3]: # Just run this cell.
pop_restaurants = Table.read_table('restaurant_inspection_scores.csv').drop(5,6)
pop_restaurants
```

```
[3]: Restaurant Name | Zip Code | Inspection Date | Score | Address
6M Grocery         | 78652    | 01/17/2014      | 90    | 805 W FM 1626 RD
AUSTIN, TX 78652
6M Grocery         | 78652    | 04/27/2015      | 93    | 805 W FM 1626 RD
AUSTIN, TX 78652
6M Grocery         | 78652    | 05/02/2016      | 88    | 805 W FM 1626 RD
AUSTIN, TX 78652
6M Grocery         | 78652    | 07/25/2014      | 100   | 805 W FM 1626 RD
AUSTIN, TX 78652
6M Grocery         | 78652    | 10/21/2015      | 87    | 805 W FM 1626 RD
AUSTIN, TX 78652
6M Grocery         | 78652    | 12/15/2014      | 93    | 805 W FM 1626 RD
```

```

AUSTIN, TX 78652
7 Eleven #36575 | 78660 | 01/25/2016 | 92 | 15829 N IH 35 SVRD NB
AUSTIN, TX 78660
7 Eleven #36575 | 78660 | 03/05/2015 | 86 | 15829 N IH 35 SVRD NB
AUSTIN, TX 78660
7 Eleven #36575 | 78660 | 03/14/2014 | 93 | 15829 N IH 35 SVRD NB
AUSTIN, TX 78660
7 Eleven #36575 | 78660 | 07/27/2015 | 97 | 15829 N IH 35 SVRD NB
AUSTIN, TX 78660
... (24357 rows omitted)


```

Question 1.1 Plot a histogram of the scores in the cell below.

```

[4]: # Write your code here.
pop_restaurants.hist("Score")

```



This is the **population mean**:

```

[5]: pop_mean = np.mean(pop_restaurants.column(3))
pop_mean

```

```

[5]: 91.40706693478886

```

Often it is impossible to find complete datasets like this. Imagine we instead had access only to

a random sample of 100 restaurant inspections, called `restaurant_sample`. That table is created below. We are interested in using this sample to estimate the population mean.

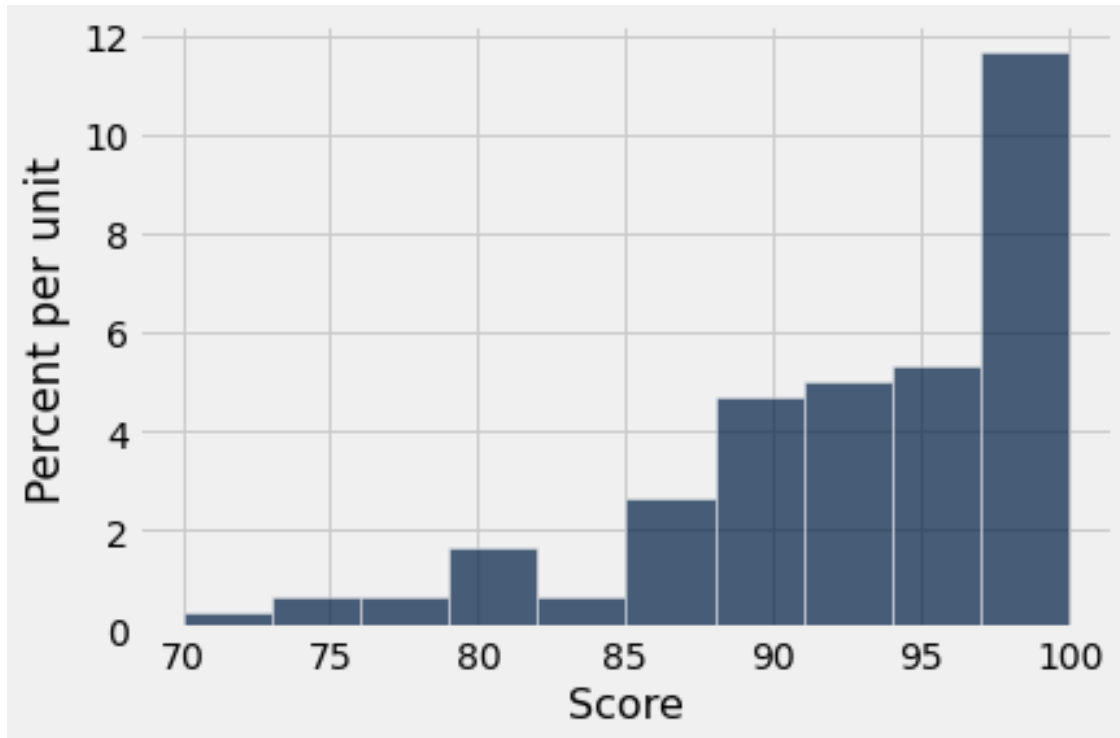
```
[6]: restaurant_sample = pop_restaurants.sample(100, with_replacement=False)
restaurant_sample
```

```
[6]: Restaurant Name | Zip Code | Inspection Date | Score |
Address
Sustainable Food Center | 78702 | 10/30/2014 | 100 |
2921 E 17TH ST Bldg C
AUSTIN, TX 78702
(30.279348, -97.7 ...
Uncle Billy's Brew & Que | 78704 | 03/17/2015 | 85 |
1530 BARTON SPRINGS RD
AUSTIN, TX 78704
(30.261942, -97. ...
Star of Texas | 78705 | 12/16/2015 | 99 |
611 W 22ND ST
AUSTIN, TX 78705
(30.285196, -97.744645)
St. Andrew's Episcopal School-Cafeteria | 78735 | 05/26/2015 | 96 |
5901 SOUTHWEST PKWY
AUSTIN, TX 78735
(30.247106, -97.848037)
"CB's" at Stubb's Bar-B-Q | 78701 | 12/02/2014 | 97 |
809 RED RIVER ST
AUSTIN, TX 78701
(30.269076, -97.736231)
Salam International Mart & Cafe | 78753 | 10/11/2016 | 70 |
10009 N LAMAR BLVD Unit B
AUSTIN, TX 78753
(30.370063, - ...
Jack in the Box #865 | 78752 | 07/14/2016 | 81 |
6419 AIRPORT BLVD
AUSTIN, TX 78752
(30.329371, -97.715771)
Whole Foods - Meat Market | 78703 | 03/12/2015 | 97 |
525 N LAMAR BLVD
AUSTIN, TX 78703
(30.270991, -97.754251)
Lorraine Grandma Camacho Ctr | 78702 | 01/09/2015 | 100 |
34 ROBERT T MARTINEZ JR ST
AUSTIN, TX 78702
(30.250227, ...
Family Dollar Store # 10028 | 78725 | 10/29/2014 | 94 |
14001 FM 969 RD
AUSTIN, TX 78725
```

... (90 rows omitted)

Question 1.2 Plot a histogram of the **sample** scores in the cell below.

```
[7]: # Write your code here:  
restaurant_sample.hist("Score")
```



This is the **sample mean**:

```
[8]: sample_mean = np.mean(restaurant_sample.column(3))  
sample_mean
```

```
[8]: 92.32
```

Question 1.3 Complete the function **bootstrap_scores** below. It should take no arguments. It should simulate drawing 5000 resamples from **restaurant_sample** and computing the mean restaurant score in each resample. It should return an array of those 5000 resample means.

```
[9]: def bootstrap_scores():  
    resampled_means = []  
    for i in range(5000):  
        resampled_mean = restaurant_sample.sample()  
        resampled_means.append(resampled_mean["Score"].mean())  
    return np.array(resampled_means)
```

```
resampled_means = bootstrap_scores()
resampled_means
```


```
[9]: array([92.84, 92.49, 92.22, ..., 92.67, 91.86, 92.12])
```

```
[10]: check('tests/q1_3.py')
```

```
[10]: <gofer.ok.OKTestsResult at 0x7fd319d4a950>
```

Take a look at the histogram of the **resampled means**.

```
[11]: Table().with_column('Resampled Means', resampled_means).hist()
```



Question 1.4 Compute a 95% confidence interval for the average restaurant score using the array `resampled_means`.

```
[12]: lower_bound = percentile(2.5, resampled_means)
upper_bound = percentile(97.5, resampled_means)
print("95% confidence interval for the average restaurant score, computed by_\n
↳bootstrapping:\n(", lower_bound, ", ", upper_bound, ")")
```

```
95% confidence interval for the average restaurant score, computed by
bootstrapping:
( 90.93 , 93.62 )
```

Question 1.5 Does the distribution of the resampled mean scores look normally distributed? State “yes” or “no” and describe in one sentence why you would expect that result.

Yes, because central limit theorem

Question 1.6 Does the distribution of the **sampled scores** look normally distributed? State “yes” or “no” and describe in one sentence why you should expect this result.

Hint: Remember that we are no longer talking about the resampled means!

Small sample + single sample so CLT doesn’t really happen.

For the last question, you’ll need to recall two facts. 1. If a group of numbers has a normal distribution, around 95% of them lie within 2 standard deviations of their mean. 2. The Central Limit Theorem tells us the quantitative relationship between the following: * The standard deviation of an array of numbers. * The standard deviation of an array of means of samples taken from those numbers.

Question 1.7 Without referencing the array `resampled_means` or performing any new simulations, calculate an interval around the `sample_mean` that covers approximately 95% of the numbers in the `resampled_means` array. **You may use the following values to compute your result, but you should not perform additional resampling** - think about how you can use the Central Limit Theorem (CLT) to accomplish this.

```
[13]: sample_mean = np.mean(restaurant_sample.column(3))
      sample_sd = np.std(restaurant_sample.column(3))
      sample_size = restaurant_sample.num_rows

      lower_bound_normal = sample_mean - 2*sample_sd
      upper_bound_normal = sample_mean + 2*sample_sd
      print("95% confidence interval for the average restaurant score, computed by a_
↪normal approximation:\n(", lower_bound_normal, ", ", upper_bound_normal, ")")
```

```
95% confidence interval for the average restaurant score, computed by a normal
approximation:
( 78.6675130470672 , 105.97248695293278 )
```

This confidence interval should look very similar to the one you computed in **Question 1.4**.

1.2 2. Testing the Central Limit Theorem

The Central Limit Theorem tells us that the probability distribution of the **sum** or **average** of a large random sample drawn with replacement will be roughly normal, *regardless of the distribution of the population from which the sample is drawn*.

That’s a pretty big claim, but the theorem doesn’t stop there. It further states that the standard deviation of this normal distribution is given by:

$$\frac{\text{sd of the original distribution}}{\sqrt{\text{sample size}}}$$

In other words, suppose we start with *any distribution* that has standard deviation x , take a sample

of size n (where n is a large number) from that distribution with replacement, and compute the **mean** of that sample. If we repeat this procedure many times, then those sample means will have a normal distribution with standard deviation $\frac{\sigma}{\sqrt{n}}$.

That's an even bigger claim than the first one! The proof of the theorem is beyond the scope of this class, but in this exercise, we will be exploring some data to see the CLT in action.

Question 2.1 The CLT only applies when sample sizes are “sufficiently large.” This isn't a very precise statement. Is 10 large? How about 50? The truth is that it depends both on the original population distribution and just how “normal” you want the result to look. Let's use a simulation to get a feel for how the distribution of the sample mean changes as sample size goes up.

Consider a coin flip. If we say **Heads** is 1 and **Tails** is 0, then there's a 50% chance of getting a 1 and a 50% chance of getting a 0, which definitely doesn't match our definition of a normal distribution. The average of several coin tosses, where Heads is 1 and Tails is 0, is equal to the proportion of heads in those coin tosses (which is equivalent to the mean value of the coin tosses), so the CLT should hold **true** if we compute the sample proportion of heads many times.

Write a function called `sample_size_n` that takes in a sample size n . It should return an array that contains 5000 sample proportions of heads, each from n coin flips.

```
[14]: sample_proportions(1, [.5, .5])
```

```
[14]: array([0., 1.])
```

```
[15]: def sample_size_n(n):
      coin_proportions = make_array(.5, .5) # our coin is fair
      heads_proportions = make_array()
      for i in np.arange(5000):
          simulated_proportions = sample_proportions(n, coin_proportions)
          prop_heads = simulated_proportions[0]
          heads_proportions = np.append(heads_proportions, prop_heads)
      return heads_proportions

      sample_size_n(5)
```

```
[15]: array([0.6, 0.8, 0.4, ..., 0.6, 0.4, 0.6])
```

The code below will use the function you just defined to plot the empirical distribution of the sample mean for various sample sizes. Drag the slider or click on the number to the right to type in a sample size of your choice. The x- and y-scales are kept the same to facilitate comparisons. Notice the shape of the graph as the sample size increases and decreases.

```
[16]: # Just run this cell
      from ipywidgets import interact

      def outer(f):
          def graph(x):
              bins = np.arange(-0.01, 1.05, 0.02)
```



```

        sample_props = f(x)
        Table().with_column('Sample Size: {}'.format(x), sample_props).
↪hist(bins=bins)
        plt.ylim(0, 30)
        print('Sample SD:', np.std(sample_props))
        plt.show()
        return graph

interact(outer(sample_size_n), x=(0, 400, 1), continuous_update=False);

# Min sample size is 0, max is 400
# The graph will refresh a few times when you drag the slider around

```

```

interactive(children=(IntSlider(value=200, description='x', max=400), Output()),
↪_dom_classes=('widget-interac...

```

You can see that even the means of samples of 10 items follow a roughly bell-shaped distribution. A sample of 50 items looks quite bell-shaped.

Question 2.2 In the plot for a sample size of 10, why are the bars spaced at intervals of .1, with gaps in between?

The bins?

Now we will test the second claim of the CLT: That the SD of the sample mean is the SD of the original distribution, divided by the square root of the sample size.

We have imported the flight delay data and computed its standard deviation for you.

```

[17]: united = Table.read_table('united_summer2015.csv')
        united_std = np.std(united.column('Delay'))
        united_std

```

```

[17]: 39.480199851609314

```

Question 2.3 Write a function called `empirical_sample_mean_sd` that takes a sample size `n` as its argument. The function should simulate 500 samples with replacement of size `n` from the flight delays dataset, and it should return the standard deviation of the **means of those 500 samples**.

Hint: This function will be similar to the `sample_size_n` function you wrote earlier.

```

[18]: def empirical_sample_mean_sd(n):
        sample_means = make_array()
        for i in np.arange(500):
            sample = united.sample(n)
            sample_mean = sample["Delay"].mean()
            sample_means = np.append(sample_means, sample_mean)
        return np.std(sample_means)

        empirical_sample_mean_sd(10)

```

```
[18]: 12.95666237732542
```

```
[19]: check('tests/q2_3.py')
```

```
[19]: <gofer.ok.OKTestsResult at 0x7fd319510590>
```

2 I have a big problem with appending to an array because arrays should be a fixed size. I'd be much more comfortable if we were to append to a list and cast it to an array at the end (which usually isn't even necessary since it's put into a table)

Question 2.4 Now, write a function called `predict_sample_mean_sd` to find the predicted value of the standard deviation of means according to the relationship between the standard deviation of the sample mean and sample size that is discussed [here](#) in the textbook. It takes a sample size `n` (a number) as its argument. It returns the predicted value of the standard deviation of the mean delay time for samples of size `n` from the flight delays (represented in the table `united`).

```
[20]: def predict_sample_mean_sd(n):  
      return united_std / np.sqrt(n)  
  
predict_sample_mean_sd(10)
```


```
[20]: 12.484735400972708
```

```
[21]: check('tests/q2_4.py')
```

```
[21]: <gofer.ok.OKTestsResult at 0x7fd3191b3ad0>
```

The cell below will plot the predicted and empirical SDs for the delay data for various sample sizes.

```
[22]: sd_table = Table().with_column('Sample Size', np.arange(1,101))  
predicted = sd_table.apply(predict_sample_mean_sd, 'Sample Size')  
empirical = sd_table.apply(empirical_sample_mean_sd, 'Sample Size')  
sd_table = sd_table.with_columns('Predicted SD', predicted, 'Empirical SD',  
    ↪ empirical)  
sd_table.scatter('Sample Size')
```



Question 2.5 Do our predicted and empirical values match? Why is this the case?

Hint: Are there any laws that we learned about in class that might help explain this?

CLT says that SD of sample means is $\text{pop_SD} / \sqrt{\text{sample_size}}$

2.1 3. Polling and the Normal Distribution

Question 3.1 Michelle is a statistical consultant, and she works for a group that supports Proposition 68 (which would mandate labeling of all horizontal or vertical axes), called Yes on 68. They want to know how many Californians will vote for the proposition.

Michelle polls a uniform random sample of all California voters, and she finds that 210 of the 400 sampled voters will vote in favor of the proposition. Fill in the code below to form a table with 3 columns: the first two columns should be identical to `sample`. The third column should be named `Proportion` and have the proportion of total voters that chose each option.

```
[23]: sample = Table().with_columns(
    "Vote", make_array("Yes", "No"),
    "Count", make_array(210, 190))
sample_size = 210 + 190
sample_with_proportions = sample.with_column("Proportion", [210 / sample_size,
↪ 190 / sample_size])
sample_with_proportions
```


```
[23]: Vote | Count | Proportion
      Yes | 210 | 0.525
      No | 190 | 0.475
```

```
[24]: check('tests/q3_1.py')
```

```
[24]: <gofer.ok.OKTestsResult at 0x7fd3190be090>
```

Question 3.2 She then wants to use 10,000 bootstrap resamples to compute a confidence interval for the proportion of all California voters who will vote Yes. Fill in the next cell to simulate an empirical distribution of Yes proportions with 10,000 resamples. In other words, use bootstrap resampling to simulate 10,000 election outcomes, and populate `resample_yes_proportions` with the yes proportion of each bootstrap resample. Then, visualize `resample_yes_proportions` with a histogram. You should see a bell shaped curve centered near the proportion of Yes in the original sample.

```
[25]: resample_yes_proportions = make_array()
      for i in np.arange(10000):
          resample = sample_proportions(sample_size, [210 / sample_size, 190 /
          ↪sample_size])
          resample_yes_proportions = np.append(resample_yes_proportions, resample[0])
      Table().with_column("Resample Yes proportion", resample_yes_proportions).
      ↪hist(bins=np.arange(.2, .8, .01))
```



```
[26]: check('tests/q3_2.py')
```

[26]: <gofer.ok.OKTestsResult at 0x7fd31905a2d0>

Question 3.3 Why does the Central Limit Theorem (CLT) apply in this situation, and how does it explain the distribution we see above?

Curve is approximately normal because we took enough samples from a distribution

In a population whose members are 0 and 1, there is a simple formula for the standard deviation of that population:

$$\text{standard deviation} = \sqrt{(\text{proportion of 0s}) \times (\text{proportion of 1s})}$$

(Figuring out this formula, starting from the definition of the standard deviation is an fun exercise for those who enjoy Algebra.)

Question 3.4 Using only the CLT and the numbers of Yes and No voters in our sample of 400, compute (*algebraically*) a number `approximate_sd` that's the predicted standard deviation of the array `resample_yes_proportions` according to the Central Limit Theorem. **Do not access the data in `resample_yes_proportions` in any way.** Remember that a predicted standard deviation of the sample means can be computed from the population SD and the size of the sample.

Also, remember that if we do not know the population SD, we can use the sample SD as a reasonable approximation in its place.

```
[27]: approximate_sd = np.sqrt((210 / sample_size) * (190 / sample_size)) / np.  
      ↪sqrt(sample_size)  
      approximate_sd
```

[27]: 0.024968730444297725

```
[28]: check('tests/q3_4.py')
```

[28]: <gofer.ok.OKTestsResult at 0x7fd318f69390>

Question 3.5

Compute the SD of the array `resample_yes_proportions` which will act as an approximation to the true SD of the possible sample proportions. This will help verify whether your answer to Question 3.2 is approximately correct.

```
[29]: exact_sd = np.std(resample_yes_proportions)  
      exact_sd
```

[29]: 0.025050897886852285

```
[30]: check('tests/q3_5.py')
```

[30]: <gofer.ok.OKTestsResult at 0x7fd318ef51d0>

Question 3.6 Again, without accessing `resample_yes_proportions` in any way, compute an approximate 95% confidence interval for the proportion of Yes voters in California.

The cell below draws your interval as a red bar below the histogram of `resample_yes_proportions`; use that to verify that your answer looks right.

Hint: How many SDs corresponds to 95% of the distribution promised by the CLT? Recall the discussion in the textbook here.


```
[31]: lower_limit = (210/sample_size) - 2*approximate_sd
      upper_limit = (210/sample_size) + 2*approximate_sd
      print('lower:', lower_limit, 'upper:', upper_limit)
```

```
lower: 0.47506253911140456 upper: 0.5749374608885954
```

```
[32]: check('tests/q3_6.py')
```

```
[32]: <gofer.ok.OKTestsResult at 0x7fd318ef7350>
```

```
[33]: # Run this cell to plot your confidence interval.
      Table().with_column("Resample Yes proportion", resample_yes_proportions).
        ↪ hist(bins=np.arange(.2, .8, .01))
      plt.plot(make_array(lower_limit, upper_limit), make_array(0, 0), c='r', lw=10);
```



Your confidence interval should overlap the number 0.5. That means we can't be very sure whether Proposition 68 is winning, even though the sample Yes proportion is a bit above 0.5.

The Yes on 68 campaign really needs to know whether they're winning. It's impossible to be absolutely sure without polling the whole population, but they'd be okay if the standard deviation of the sample mean was only 0.005. They ask Michelle to run a new poll with a sample size that's large enough to achieve that. (Polling is expensive, so the sample also shouldn't be bigger than necessary.)

Michelle consults Chapter 14 of your textbook. Instead of making the conservative assumption that the population standard deviation is 0.5 (coding Yes voters as 1 and No voters as 0), she decides to assume that it's equal to the standard deviation of the sample:

$$\sqrt{(\text{Yes proportion in the sample}) \times (\text{No proportion in the sample})}.$$

Under that assumption, Michelle decides that a sample of 9,975 would suffice.

Question 3.7 Does Michelle's sample size achieve the desired standard deviation of sample means? What SD would you achieve with a smaller sample size? A higher sample size? To explore this, first compute the SD of sample means obtained by using Michelle's sample size.

```
[34]: estimated_population_sd = np.sqrt(210/400 * 190/400)
michelle_sample_size = 9975
michelle_sample_mean_sd = estimated_population_sd / np.
    ↪sqrt(michelle_sample_size)
print("With Michelle's sample size, you would predict a sample mean SD of %f."
    ↪% michelle_sample_mean_sd)
```

With Michelle's sample size, you would predict a sample mean SD of 0.005000.

Then, compute the SD of sample means that you would get from a smaller sample size. Ideally, you should pick a number that is significantly smaller, but any sample size smaller than Michelle's will do.

```
[35]: smaller_sample_size = 9000
smaller_sample_mean_sd = estimated_population_sd / np.sqrt(smaller_sample_size)
print("With this smaller sample size, you would predict a sample mean SD of %f"
    ↪% smaller_sample_mean_sd)
```

With this smaller sample size, you would predict a sample mean SD of 0.005264

Finally, compute the SD of sample means that you would get from a larger sample size. Here, a number that is significantly larger would make any difference more obvious, but any sample size larger than Michelle's will do.

```
[36]: larger_sample_size = 20000
larger_sample_mean_sd = estimated_population_sd / np.sqrt(larger_sample_size)
print("With this larger sample size, you would predict a sample mean SD of %f"
    ↪% larger_sample_mean_sd)
```

With this larger sample size, you would predict a sample mean SD of 0.003531

```
[37]: check('tests/q3_7.py')
```

```
[37]: <gofer.ok.OKTestsResult at 0x7fd318e17bd0>
```

Question 3.8 Based off of this, was Michelle’s sample size approximately the minimum sufficient sample, given her assumption that the sample SD is the same as the population SD? Assign `min_sufficient` to `True` if this 9975 was indeed approximately the minimum sufficient sample, and `False` if it wasn’t.

```
[38]: min_sufficient = True
      min_sufficient
```

```
[38]: True
```

```
[39]: check('tests/q3_8.py')
```

```
[39]: <gofer.ok.OKTestsResult at 0x7fd318e18a50>
```

2.2 4. Submission

Once you’re finished, select “Save and Checkpoint” in the File menu. Your name and course section number should be in the first and last cell of the assignment. Be sure you have run all cells with code and that the output from that is showing.

Double check that you have completed all of the free response questions as the auto-grader does NOT check that and YOU are responsible for knowing those questions are there and completing them as part of the grade for this homework. When ready, click “Print Preview” in the File menu. Print a copy from there in pdf format. (This means you right click and choose print and choose “save as pdf” from your printer options.) You will need to submit the pdf in Canvas by the deadline.

The gopher grader output and/or output from your coding are essential to helping your instructor grade your work correctly and in a timely manner.

Files submitted that are missing the required output will lose some to all points so double check your pdf before submitting.

```
[40]: # For your convenience, you can run this cell to run all the tests at once!
import glob
from gofer.ok import grade_notebook
if not globals().get('__GOFER_GRADER__', False):
    display(grade_notebook('hw09.ipynb', sorted(glob.glob('tests/q*.py'))))
```

95% confidence interval for the average restaurant score, computed by bootstrapping:

(88.8 , 91.86)

95% confidence interval for the average restaurant score, computed by a normal approximation:

(74.616278253382 , 106.08372174661798)

lower: 0.47506253911140456 upper: 0.5749374608885954

With Michelle's sample size, you would predict a sample mean SD of 0.005000.

With this smaller sample size, you would predict a sample mean SD of 0.005264

With this larger sample size, you would predict a sample mean SD of 0.003531

```
['tests/q1_3.py', 'tests/q2_3.py', 'tests/q2_4.py', 'tests/q3_1.py',  
'tests/q3_2.py', 'tests/q3_4.py', 'tests/q3_5.py', 'tests/q3_6.py',  
'tests/q3_7.py', 'tests/q3_8.py']
```

Question 1:

<gofer.ok.OKTestsResult at 0x7fd31894e0d0>

Question 2:

<gofer.ok.OKTestsResult at 0x7fd319127a90>

Question 3:

<gofer.ok.OKTestsResult at 0x7fd3193faa50>

Question 4:

<gofer.ok.OKTestsResult at 0x7fd318bf74d0>

Question 5:

<gofer.ok.OKTestsResult at 0x7fd3194b82d0>

Question 6:

<gofer.ok.OKTestsResult at 0x7fd319455ad0>

Question 7:

<gofer.ok.OKTestsResult at 0x7fd319455c10>

Question 8:

<gofer.ok.OKTestsResult at 0x7fd3193965d0>



Question 9:



<gofer.ok.OKTestsResult at 0x7fd319598610>



Question 10:

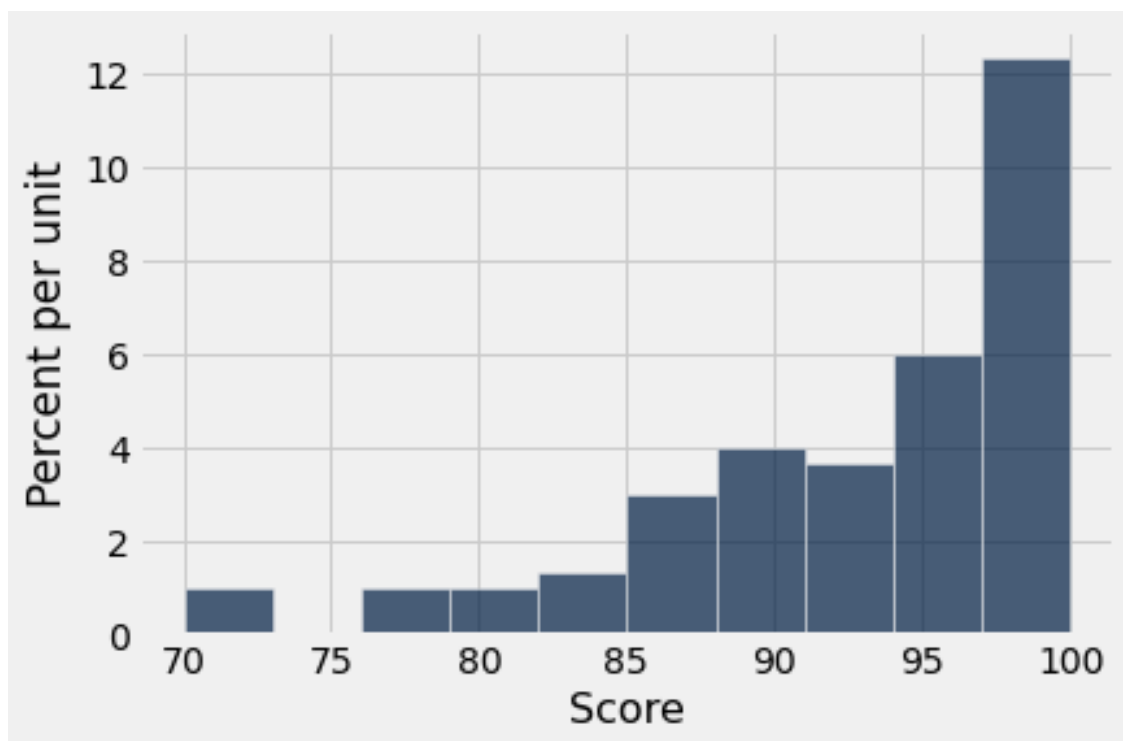

<gofer.ok.OKTestsResult at 0x7fd319396e50>



1.0















Name: Allan Gongora

Section: 0131