

hw11

June 2, 2022

Name: Allan Gongora

Section: 0131

1 Homework 11: Regression Inference

Please complete this notebook by filling in the cells provided. Before you begin, execute the following cell to load the provided tests.

```
[1]: pip install gofer-grader
```

Collecting gofer-grader

Using cached gofer_grader-1.1.0-py3-none-any.whl (9.9 kB)

Requirement already satisfied: tornado in /opt/conda/lib/python3.7/site-packages (from gofer-grader) (6.1)

Requirement already satisfied: pygments in /opt/conda/lib/python3.7/site-packages (from gofer-grader) (2.11.2)

Requirement already satisfied: jinja2 in /opt/conda/lib/python3.7/site-packages (from gofer-grader) (3.0.3)

Requirement already satisfied: MarkupSafe>=2.0 in /opt/conda/lib/python3.7/site-packages (from jinja2->gofer-grader) (2.0.1)

Installing collected packages: gofer-grader

Successfully installed gofer-grader-1.1.0

Note: you may need to restart the kernel to use updated packages.

```
[2]: # Don't change this cell; just run it.
```

```
import numpy as np
from datascience import *

# These lines do some fancy plotting magic.
import matplotlib
%matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
import warnings
warnings.simplefilter('ignore', FutureWarning)
from matplotlib import patches
```

```

from ipywidgets import interact, interactive, fixed
import ipywidgets as widgets

# These lines load the tests.

from gofer.ok import check

```

Recommended Reading: * [Inference for Regression](#) 1) For all problems that you must write explanations and sentences for, you **must** provide your answer in the designated space. This can include: A) Sentence responses to questions that ask for an explanation B) Numeric responses to multiple choice questions C) Programming code

- 2) Moreover, throughout this homework and all future ones, please be sure to not re-assign variables throughout the notebook! For example, if you use `max_temperature` in your answer to one question, do not reassign it later on. Otherwise, you will fail tests that you thought you were passing previously!

Once you're finished, select "Save and Checkpoint" in the File menu. Your name and course section number should be in the first and last cell of the assignment. Be sure you have run all cells with code and that the output from that is showing. Then click "Print Preview" in the File menu. Print a copy from there in pdf format. (This means you right click and choose print and choose "save as pdf" from your printer options.) You will need to submit the pdf in Canvas by the deadline.

The gopher grader output and/or output from your coding are essential to helping your instructor grade your work correctly and in a timely manner.

Files submitted that are missing the required output will lose some to all points so double check your pdf before submitting.

1.1 1. Visual Diagnostics for Linear Regression

1.1.1 Regression Model Diagnostics

Linear regression isn't always the best way to describe the relationship between two variables. We'd like to develop techniques that will help us decide whether or not to use a linear model to predict one variable based on another.

We will use the insight that if a regression fits a set of points well, then the residuals from that regression line will show no pattern when plotted against the predictor variable.

The table below contains information about crime rates and median home values in suburbs of Boston. We will attempt to use linear regression to predict median home values in terms of crime rate.

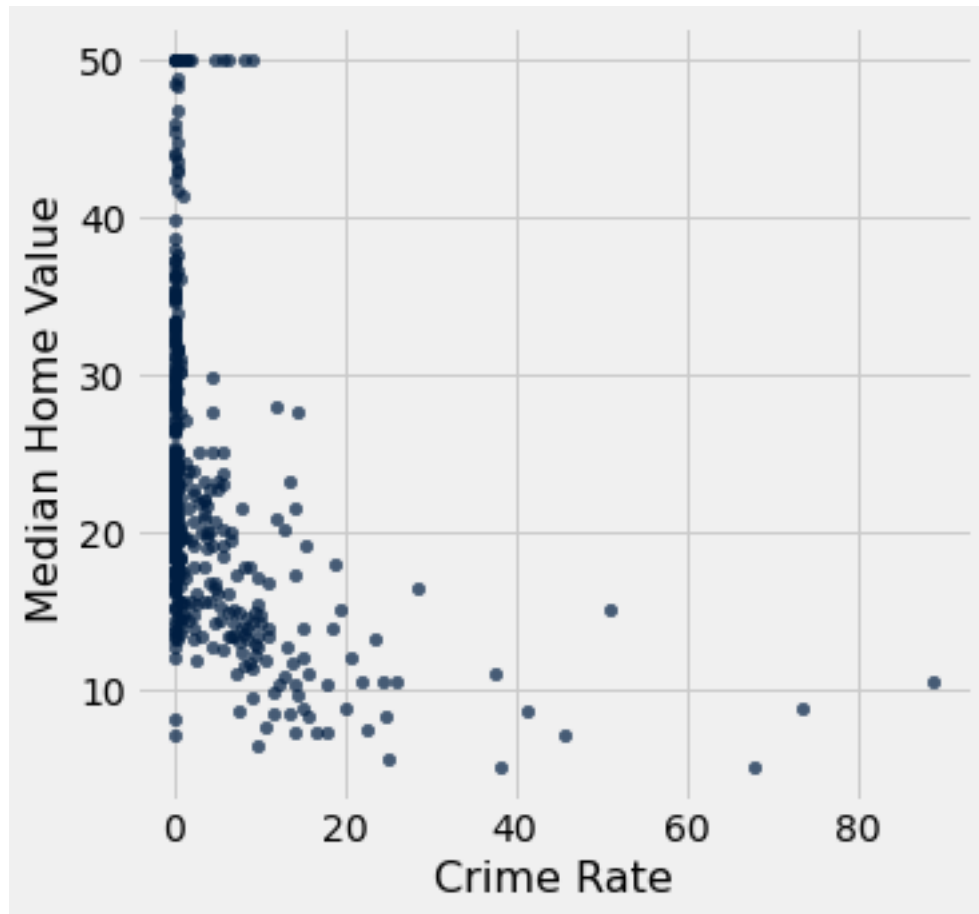
About the Dataset All data are from 1970. Crime rates are per capita per year; home values are in thousands of dollars. The crime data comes from the FBI, and home values are from the US Census Bureau.

Run the next cell to load the data and see a scatter plot.

```

[3]: boston = Table.read_table('boston_housing.csv')
      boston.scatter('Crime Rate')

```



Question 1.1 Write a function called `residuals`. It should take a single argument, a table. It should first compute the slope and intercept of the regression line that predicts the second column of that table (accessible as `tbl.column(1)`) using the first column (`tbl.column(0)`). The function should return an array containing the *residuals* for that regression line. Recall that residuals are given by:

$$\text{residual} = \text{observed value} - \text{regression estimate}$$

Hint: If your code is getting long, think about how you can split the problem up into multiple smaller, simpler functions.

```
[4]: def errors(t, slope, intercept):
      return t[1] - (t[0]*slope + intercept)

      def fit_line(tbl):
          # Your code may need more than 1 line below here.
          def mse(m, b):
              return np.mean(errors(tbl, m, b)**2)
          return minimize(mse)
```

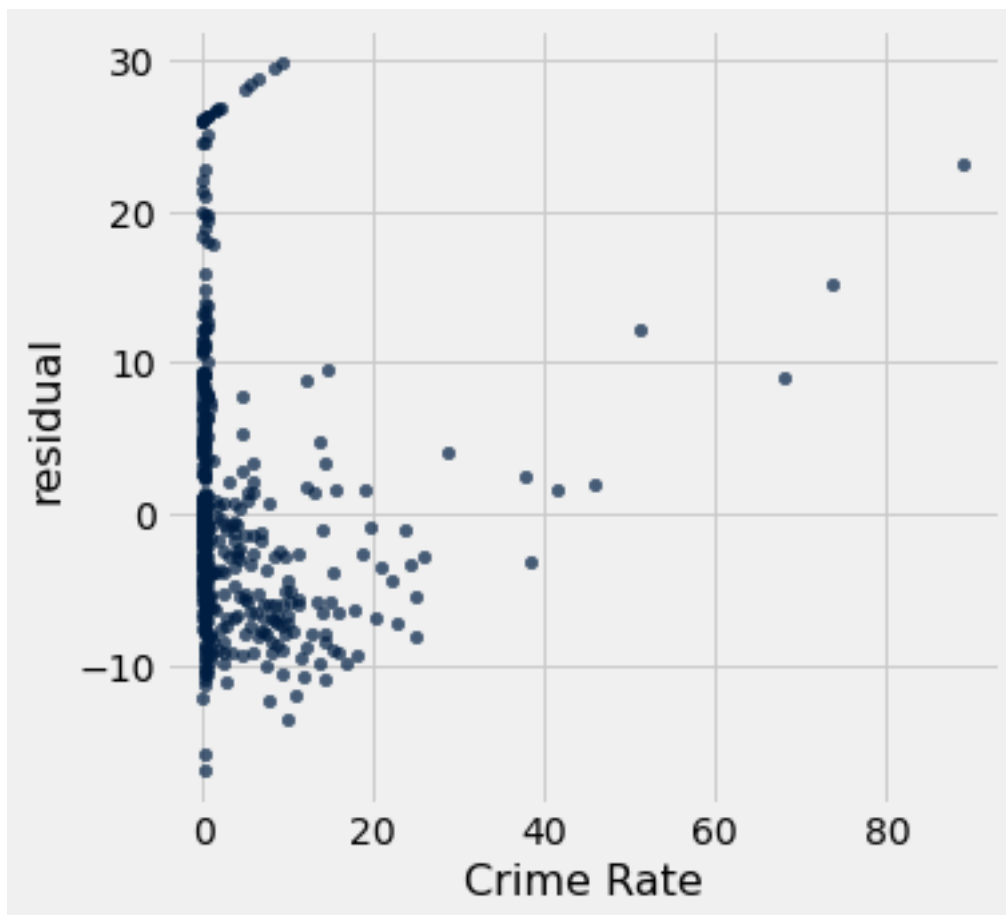
```
def residuals(tbl):  
    slope, intercept = fit_line(tbl)  
    return tbl.column(1) - (tbl.column(0)*slope + intercept)
```

```
[5]: check('tests/q1_1.py')
```

```
[5]: <gofer.ok.OKTestsResult at 0x7f93569ab390>
```

Question 1.2 Make a scatter plot of the residuals for the Boston housing dataset against crime rate. Crime rate should be on the horizontal axis.

```
[6]: boston.with_column("residual", residuals(boston)).scatter(0, 2)
```



Question 1.3 Does the plot of residuals look roughly like a formless cloud, or is there some kind of pattern? Are they centered around 0?

Write your answer here, replacing this text.

Question 1.4 Does it seem like a linear model is appropriate for describing the relationship between crime and median home value?

Assign `linear` to `True` if a linear model is appropriate for describing the relationship, and `False` if it is not.

```
[7]: linear = True
```

```
[8]: check('tests/q1_4.py')
```

```
[8]: <gofer.ok.OKTestsResult at 0x7f935475fb50>
```

Section 15.6 of the textbook describes some mathematical facts that hold for all regression estimates, regardless of goodness of fit. One fact is that there is a relationship between the standard deviation of the residuals, the standard deviation of the response variable, and the correlation. Let's test this.

Question 1.5 Directly compute the standard deviation of the residuals from the Boston data. Then, compute the same quantity without using the residuals, using the formula described in Section 15.6 [here](#) instead.

```
[12]: def standard_units(any_numbers):  
        """Convert any array of numbers to standard units."""  
        return (any_numbers - np.average(any_numbers)) / np.std(any_numbers)  
  
def correlation(t, x, y):  
        """Return the correlation coefficient (r) of two variables."""  
        return np.mean(standard_units(t.column(x)) * standard_units(t.column(y)))
```

```
[15]: boston_residual_sd = residuals(boston).std()  
boston_residual_sd_from_formula = np.sqrt(1 - correlation(boston, 0, 1)**2) *  
    ↪ boston[1].std()  
  
print("Residual SD: {0}".format(boston_residual_sd))  
print("Residual SD from the formula: {0}".  
    ↪ format(boston_residual_sd_from_formula))
```

```
Residual SD: 8.474144781120948
```

```
Residual SD from the formula: 8.474144781120948
```

```
[16]: check('tests/q1_5.py')
```

```
[16]: <gofer.ok.OKTestsResult at 0x7f9352f19710>
```

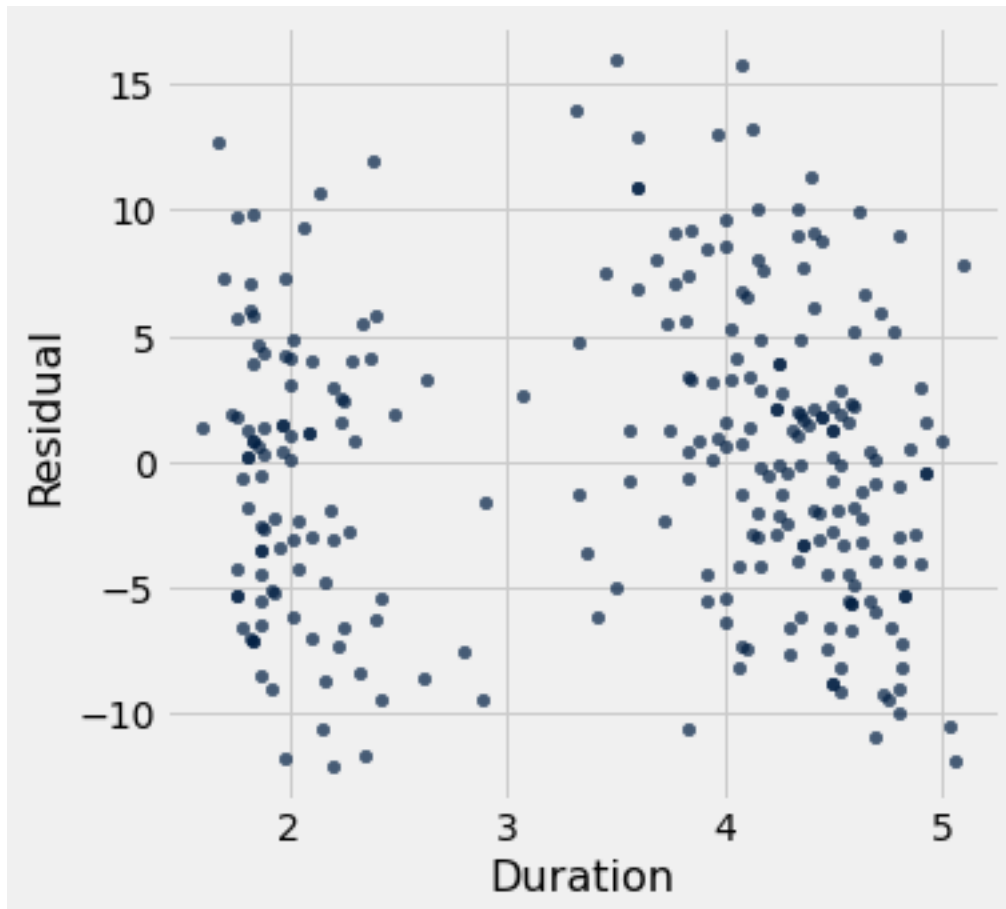
Below, we have imported a new table, the Old Faithful data.

```
[17]: old_faithful = Table.read_table('faithful.csv')  
old_faithful
```

```
[17]: duration | wait
      3.6      | 79
      1.8      | 54
      3.333    | 74
      2.283    | 62
      4.533    | 85
      2.883    | 55
      4.7      | 88
      3.6      | 85
      1.95     | 51
      4.35     | 85
      ... (262 rows omitted)
```

The following cell makes a residual plot for this new dataset.

```
[18]: Table().with_columns('Residual', residuals(old_faithful), 'Duration',
    ↪old_faithful.column('duration')).scatter('Duration')
```



Question 1.6 Repeat the procedure from Question 1.5 to find the standard deviation of the

residuals for the `old_faithful` dataset.

```
[19]: faithful_residual_sd = residuals(old_faithful).std()
      faithful_residual_sd_from_formula = np.sqrt(1 - correlation(old_faithful, 0,
↪1)**2) * old_faithful[1].std()

      print("Residual SD: {0}".format(faithful_residual_sd))
      print("Residual SD from the formula: {0}".
↪format(faithful_residual_sd_from_formula))
```

Residual SD: 5.892226635893953

Residual SD from the formula: 5.892226635893948

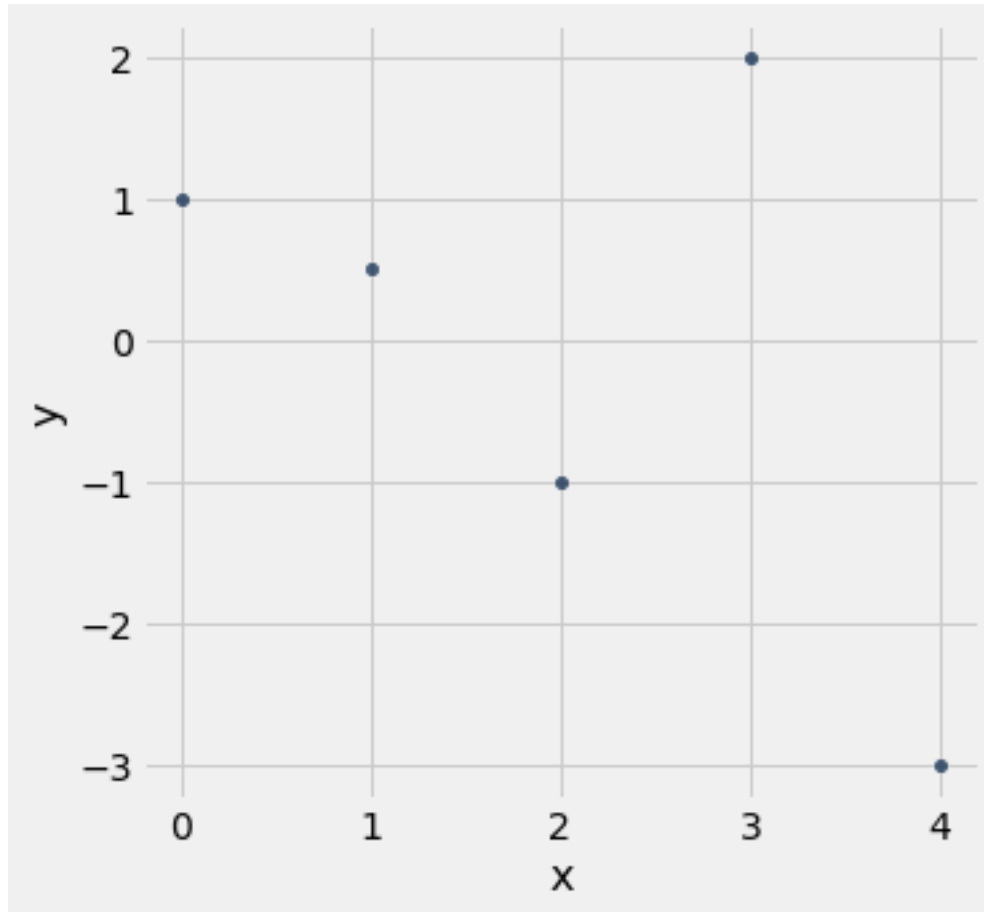
```
[20]: check('tests/q1_6.py')
```

```
[20]: <gofer.ok.OKTestsResult at 0x7f9354703650>
```

1.2 2. Finding the Least Squares Regression Line

In this exercise, you'll work with a small invented data set. Run the next cell to generate the dataset `d` and see a scatter plot.

```
[21]: d = Table().with_columns(
      'x', make_array(0, 1, 2, 3, 4),
      'y', make_array(1, .5, -1, 2, -3))
      d.scatter('x')
```



Question 2.1 (Ungraded, but you'll need the result later) Running the cell below will generate sliders that control the slope and intercept of a line through the scatter plot. When you adjust a slider, the line will move.

By moving the line around, make your best guess at the least-squares regression line. (It's okay if your line isn't exactly right, as long as it's reasonable.)

Note: Python will probably take about a second to redraw the plot each time you adjust the slider. We suggest clicking the place on the slider you want to try and waiting for the plot to be drawn; dragging the slider handle around will cause a long lag.

```
[22]: def plot_line(slope, intercept):
    plt.figure(figsize=(5,5))

    endpoints = make_array(-2, 7)
    p = plt.plot(endpoints, slope*endpoints + intercept, color='orange',
        label='Proposed line')

    plt.scatter(d.column('x'), d.column('y'), color='blue', label='Points')
```



```

plt.xlim(-4, 8)
plt.ylim(-6, 6)
plt.gca().set_aspect('equal', adjustable='box')

plt.legend(bbox_to_anchor=(1.8, .8))
plt.show()

interact(plot_line, slope=widgets.FloatSlider(min=-4, max=4, step=.1),
        ↪intercept=widgets.FloatSlider(min=-4, max=4, step=.1));

interactive(children=(FloatSlider(value=0.0, description='slope', max=4.0,
        ↪min=-4.0), FloatSlider(value=0.0, d...

```

You can probably find a reasonable guess by just eyeballing it. But remember: the least-squares regression line minimizes the mean of the squared errors made by the line for each point. Your eye might not be able to judge squared errors very well.

1.2.1 A Note on Mean and Total Squared Error

It is common to think of the least-squares line as the line with the least *mean*-squared error (or the square root of the mean-squared error), as the textbook does.

But it turns out that it doesn't matter whether you minimize the mean-squared error or the *total* squared error. You'll get the same best line in either case.

That's because the total squared error is just the mean-squared error multiplied by the number of points (`d.num_rows`). So if one line gets a better total-squared error than another line, then it also gets a better mean-squared error. In particular, the line with the smallest total squared error is also better than every other line in terms of mean-squared error. That makes it the least-squares line.

Tl; dr: Minimizing the mean-squared error minimizes the total-squared error as well.

Question 2.2 (Ungraded, but you'll need the result later) The next cell produces a more useful plot. Use it to find a line that's closer to the least-squares regression line, keeping the above note in mind.

```

[23]: def plot_line_and_errors(slope, intercept):
    plt.figure(figsize=(5,5))
    points = make_array(-2, 7)
    p = plt.plot(points, slope*points + intercept, color='orange',
    ↪label='Proposed line')
    ax = p[0].axes

    predicted_ys = slope*d.column('x') + intercept
    diffs = predicted_ys - d.column('y')
    for i in np.arange(d.num_rows):
        x = d.column('x').item(i)
        y = d.column('y').item(i)
        diff = diffs.item(i)

```

```

    if diff > 0:
        bottom_left_x = x
        bottom_left_y = y
    else:
        bottom_left_x = x + diff
        bottom_left_y = y + diff

    ax.add_patch(patches.Rectangle(make_array(bottom_left_x,
↪bottom_left_y), abs(diff), abs(diff), color='red', alpha=.3, label=('Squared_
↪error' if i == 0 else None)))
    plt.plot(make_array(x, x), make_array(y, y + diff), color='red', alpha=.
↪.6, label=('Error' if i == 0 else None))

plt.scatter(d.column('x'), d.column('y'), color='blue', label='Points')

plt.xlim(-4, 8)
plt.ylim(-6, 6)
plt.gca().set_aspect('equal', adjustable='box')

plt.legend(bbox_to_anchor=(1.8, .8))
plt.show()

interact(plot_line_and_errors, slope=widgets.FloatSlider(min=-4, max=4, step=.
↪1), intercept=widgets.FloatSlider(min=-4, max=4, step=.1));

```

```

interactive(children=(FloatSlider(value=0.0, description='slope', max=4.0,
↪min=-4.0), FloatSlider(value=0.0, d...

```

Question 2.3 Describe the visual criterion you used to find a line in Question 2.2. How did you judge whether one line was better than another?

For example, a possible (but incorrect) answer is, “I tried to make the red line for the bottom-right point as small as possible.”

I just tried to make the line as close to each point as possible

Question 2.4 We can say that a point influences the line by how much the line would move if the point was removed from the data set. Does the outlier at (3, 2) have more or less influence than any other point on the resulting best-fit line?

Assign `more_influence` to `True` if the outlier (3,2) has more influence than any other point on the best-fit line, or `False` if it does not.

```
[24]: more_influence = True
```

```
[25]: check('tests/q2_4.py')
```

```
[25]: <gofer.ok.OKTestsResult at 0x7f9352b7ee90>
```

Now, let's have Python find this line for us. When we use `minimize`, Python goes through a process similar to the one you might have used in Question 2.2.

But Python can't look at a plot that displays errors! Instead, we tell it how to find the total squared error for a line with a given slope and intercept.

Question 2.5 Define a function called `total_squared_error`. It should take two numbers as arguments:

1. The slope of some potential line
2. The intercept of some potential line

It should return the total squared error when we use that line to make predictions for the dataset `d`.

Recall that `d` has two columns: x and y .

```
[51]: def total_squared_error(slope, intercept):
      predictions = slope*d[0] + intercept
      errors = predictions - d[1]
      return sum(errors**2)
```

```
[52]: check('tests/q2_5.py')
```

```
[52]: <gofer.ok.OKTestsResult at 0x7f9352bae210>
```

Question 2.6 What is the total squared error for the line you found by “eyeballing” the errors in Question 2.1? What about Question 2.2, where you made a guess that was “aided” by a visualization of the squared error? (It's okay if the error went up, but for many students, the error will go down when using the visual aid.)

```
[35]: eyeballed_error = total_squared_error(-1.2, 2.3)
      aided_error = total_squared_error(-1.1, 2.5)
      print("Eyeballed error:", eyeballed_error, "\nAided error:", aided_error)
```

```
Eyeballed error: 13.999999999999998
```

```
Aided error: 13.799999999999999
```

```
[36]: check('tests/q2_6.py')
```

```
[36]: <gofer.ok.OKTestsResult at 0x7f9352ad3cd0>
```

Question 2.7 Use `minimize` to find the slope and intercept for the line that minimizes the total squared error. This is the definition of a least-squares regression line.

Note: `minimize` will return a single array containing the slope as the first element and intercept as the second. Read more of its documentation [here](#) or an example of its use [here](#).

```
[38]: # The staff solution used 1 line of code above here.
      slope_from_minimize, intercept_from_minimize = minimize(total_squared_error)
      print("Least-squares regression line: predicted_y =",
```

```
slope_from_minimize,  
"* x + ",  
intercept_from_minimize)
```

Least-squares regression line: $\text{predicted_y} = -0.6498374231570391 * x + 1.1988619620992735$

```
[39]: check('tests/q2_7.py')
```

```
[39]: <gofer.ok.OKTestsResult at 0x7f9352dc97d0>
```

Question 2.8 What is the total squared error for the least-squares regression line that you found?

```
[40]: best_total_squared_error = total_squared_error(slope_from_minimize,   
↪ intercept_from_minimize)  
best_total_squared_error
```

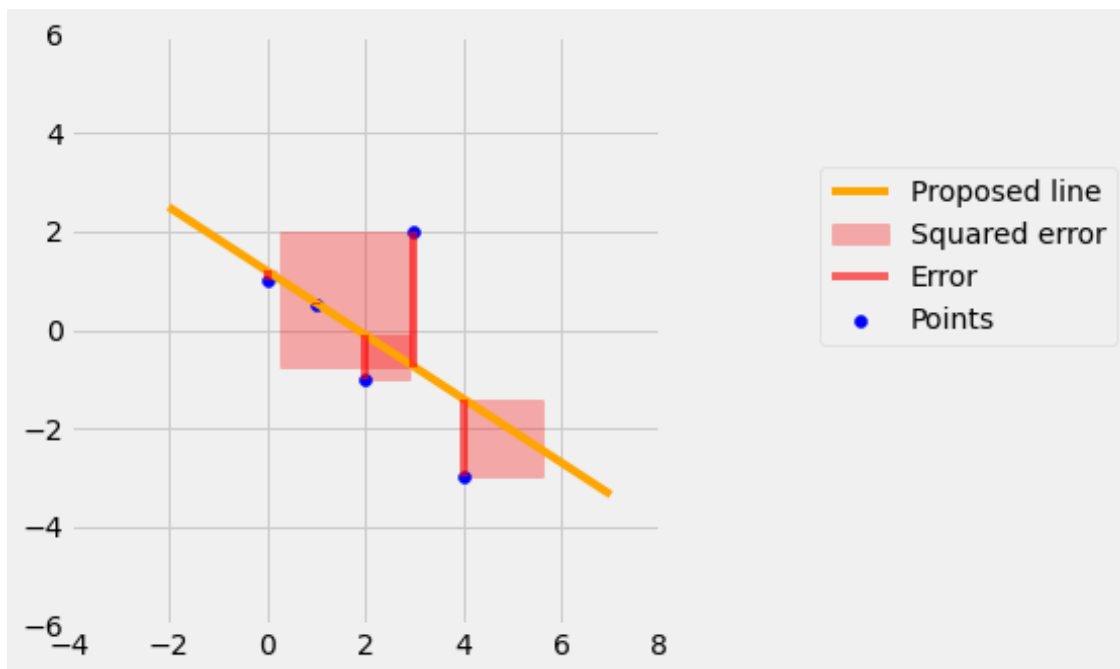
```
[40]: 10.975003568216033
```

```
[41]: check('tests/q2_8.py')
```

```
[41]: <gofer.ok.OKTestsResult at 0x7f9353008790>
```

Finally, run the following cell to plot this “best fit” line and its errors:

```
[42]: plot_line_and_errors(slope_from_minimize, intercept_from_minimize)
```



1.3 3. Quantifying Sampling Errors in Regression

Previously, in this class we've used confidence intervals to quantify uncertainty about estimates as well as to test predictions. To run a hypothesis test using a confidence interval, we use the following procedure: 1. Formulate a null hypothesis 2. Formulate an alternative hypothesis 3. Choose a test statistic and compute the observed value for the test statistic 4. Bootstrap, finding a value of the test statistic for each resample 5. Generate a 95% confidence interval from those resampled test statistics 6. Based on whether your value is in an interval, make a conclusion

Another thing we've covered recently is the use of linear regression to make predictions, using correlated variables. An example is, say, predicting the height of children based on the heights of their parents.

We can combine these two topics together in order to make even more powerful statements about our population given just a sample as before. We can use the following techniques to do so: - Bootstrapped interval for the true slope - Bootstrapped prediction interval for y (given a particular value of x)

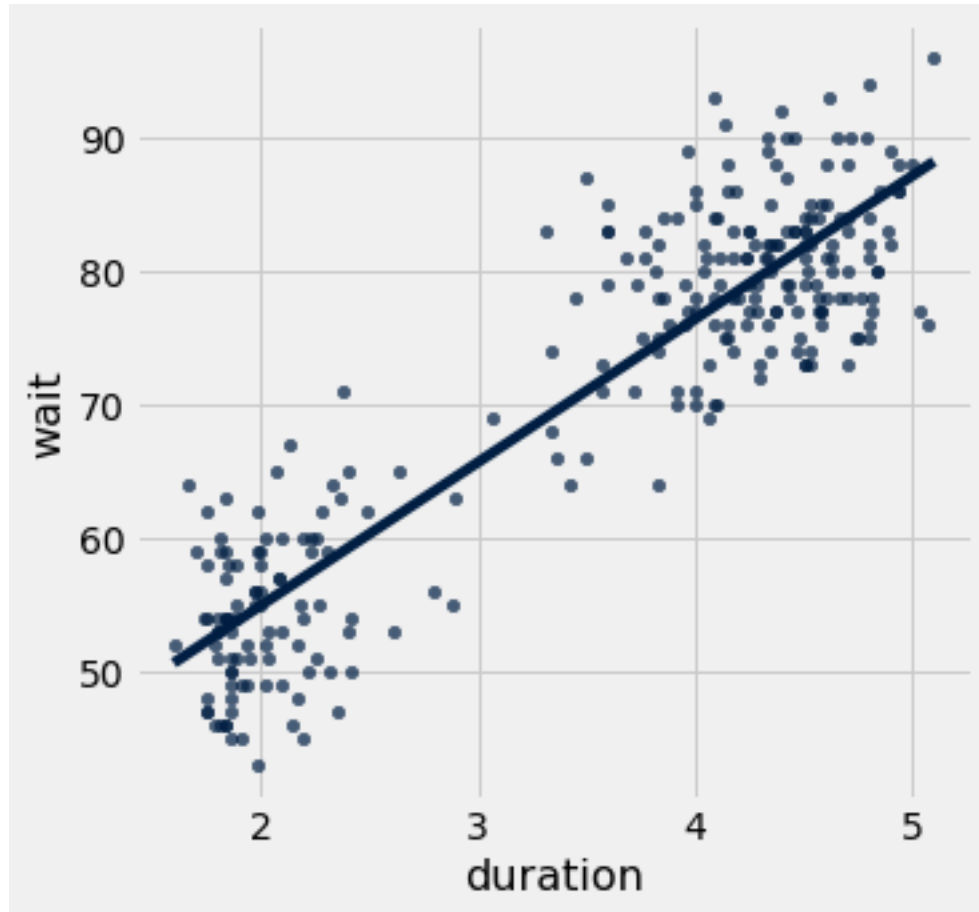
This quiz further explores these two advanced methods.

Recall the Old Faithful dataset from our lab on regression. The table contains two pieces of information about each eruption of the Old Faithful geyser in Yellowstone National Park: 1. The duration of the eruption, in minutes. 2. The time between this eruption and the next eruption (the "waiting time"), in minutes.

The dataset is plotted below along with its line of best fit.

```
[43]: faithful = Table.read_table('faithful_inference.csv')
      faithful.scatter('duration', fit_line=True)
      faithful
```

```
[43]: duration | wait
      3.6      | 79
      1.8      | 54
      3.333    | 74
      2.283    | 62
      4.533    | 85
      2.883    | 55
      4.7      | 88
      3.6      | 85
      1.95     | 51
      4.35     | 85
      ... (262 rows omitted)
```



1.3.1 Finding the Bootstrap Confidence Interval for the True Slope

Last time we looked at this dataset, we noticed the apparent linear relationship between duration and wait, and we decided to use regression to predict wait in terms of duration. However, our data is just a sample of all the eruptions that have happened at Old Faithful. As we know, relationships can appear in a sample that don't really exist in the population from which the sample was taken.

Question 3.1 Before we move forward using our linear model, we would like to know whether or not there truly exists a relationship between duration and wait time. If there is no relationship between the two, then we'd expect a correlation of 0, which would give us a slope of 0. Now, write in null and alternative hypotheses, based on your knowledge of hypothesis tests you've conducted in the past.

- **Null Hypothesis:** $\text{Corr} == 0$
- **Alternate Hypothesis:** $\text{Corr} \neq 0$

We will use the method of confidence intervals to test this hypothesis.

Question 3.2 We'll warm up by implementing some familiar functions. You may use these functions throughout this quiz. Start by defining these two functions:

1. `standard_units` should take in an array of numbers and return an array containing those numbers converted to standard units.
2. `correlation` should take in a table with 2 columns and return the correlation between these columns.

Hint: You may want to use the `standard_units` function you defined above.

```
[44]: def standard_units(any_numbers):
        """Convert any array of numbers to standard units."""
        return (any_numbers - np.average(any_numbers)) / np.std(any_numbers)

def correlation(t):
    """Return the correlation coefficient (r) of two variables."""
    return np.mean(standard_units(t.column(0)) * standard_units(t.column(1)))
```

```
[45]: check('tests/q3_2.py')
```

```
[45]: <gofer.ok.OKTestsResult at 0x7f93532c8350>
```

Question 3.3 Using the functions you just implemented, create a function called `fit_line`. It should take a table as its argument. It should return an array containing the slope and intercept of the regression line that predicts the second column in the table using the first.

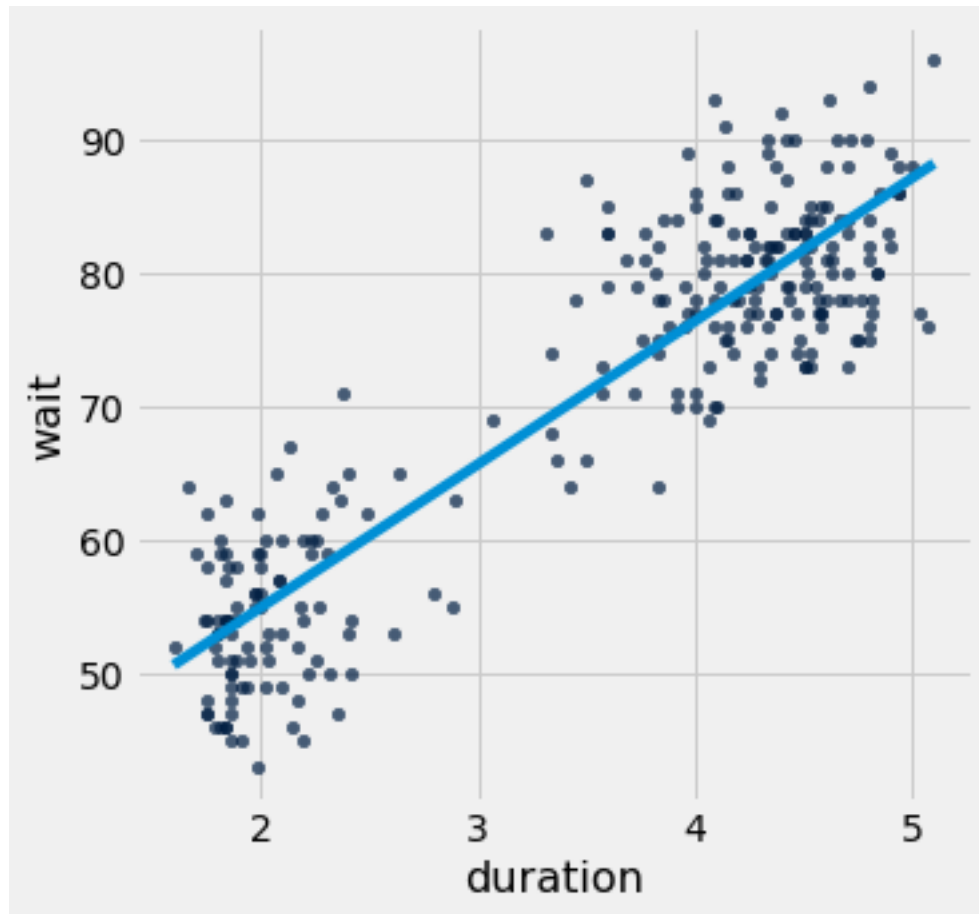
```
[50]: def fit_line(tbl):
        def mse(m, b):
            return np.mean((tbl[1] - (m*tbl[0] + b))**2)
        slope, intercept = minimize(mse)
        return make_array(slope, intercept)

# This should compute the slope and intercept of the regression
# line predicting wait time from duration in the faithful dataset.
fit_line(faithful)
```

```
[50]: array([10.72964138, 33.47439706])
```

```
[53]: # Ensure your fit_line function fits a reasonable line
# to the data in faithful, using the plot below

slope, intercept = fit_line(faithful)
faithful.scatter(0)
plt.plot([min(faithful[0]), max(faithful[0])],
         [slope*min(faithful[0])+intercept, slope*max(faithful[0])+intercept])
plt.show()
```



```
[54]: check('tests/q3_3.py')
```

```
[54]: <gofer.ok.OKTestsResult at 0x7f9353690a90>
```

Now we have all the tools we need in order to create a confidence interval quantifying our uncertainty about the true relationship between duration and wait time.

Question 3.4 Use the bootstrap to compute 1000 resamples from our dataset. For each resample, compute the slope of the best fit line. Put these slopes in an array called `resample_slopes`, giving you the empirical distribution of regression line slopes in resamples. Plot a histogram of these slopes.

```
[59]: resample_slopes = []  
      for _ in range(1000):  
          resample_slopes.append(fit_line(faithful.sample())[0])
```

Question 3.5 Use your resampled slopes to construct an approximate 95% confidence interval for the true value of the slope.


```
[60]: lower_end = np.percentile(resample_slopes, 2.5)
      upper_end = np.percentile(resample_slopes, 97.5)
      print("95% confidence interval for slope: [{:g}, {:g}]".format(lower_end,
      ↪upper_end))
```

95% confidence interval for slope: [10.1966, 11.3117]

```
[61]: check('tests/q3_5.py')
```

```
[61]: <gofer.ok.OKTestsResult at 0x7f9352fff350>
```

Question 3.6

Based on your confidence interval, would you accept or reject the null hypothesis that the true slope is 0? Why? What P-value cutoff are you using?

alpha = .05 Reject null, 95% of the data falls between 10.1966 & 11.3117 so I find it highly unlikely that the true slope is 0

1.3.2 Finding the Bootstrap Prediction Interval

Suppose we're tourists at Yellowstone, and we'd like to know how long we'll have to wait for the next Old Faithful eruption. We decide to use our regression line to make some predictions for the waiting times. But just as we're uncertain about the slope of the true regression line, we're also uncertain about the predictions we'd make based on the true regression line.

Question 3.7

Define the function `fitted_value`. It should take 2 arguments:

1. A table with 2 columns. We'll be predicting the values in the second column using the first.
2. A number, the value of the predictor variable for which we'd like to make a prediction.

Make sure to use your `fit_line` function.

```
[63]: def fitted_value(table, x):
      # The staff solution took 4 lines of code.
      m, b = fit_line(table)
      return m*x + b

      # Here's an example of how fitted_value is used. This should
      # compute the prediction for the wait time of an eruption that lasts
      # two minutes .
      two_minutes_wait = fitted_value(faithful, 2)
      two_minutes_wait
```

```
[63]: 54.933679829682376
```

```
[64]: check('tests/q3_7.py')
```

```
[64]: <gofer.ok.OKTestsResult at 0x7f93537a3450>
```

Question 3.8 The park ranger tells us that the most recent eruption lasted 6 minutes. Using your function above, assign the variable `most_recent_wait` to the predicted wait time.

```
[65]: most_recent_wait = fitted_value(faithful, 6)
      most_recent_wait
```

```
[65]: 97.85224536362261
```

```
[66]: check('tests/q3_8.py')
```

```
[66]: <gofer.ok.OKTestsResult at 0x7f9353249210>
```

Juan, a fellow tourist, raises the following objection to your prediction:

“Your prediction depends on your sample of 272 eruptions. Couldn’t your prediction have been different if you had happened to have a different sample of eruptions?”

Having read Section 16.3 of the textbook, you know just the response!

Question 3.9 Define the function `bootstrap_lines`. It should take two arguments: 1. A table with two columns. As usual, we’ll be predicting the second column using the first. 2. An integer, a number of bootstraps to run.

It should return a *table* whose first column, `"Slope"`, contains the given number of bootstrapped slopes, and whose second column, `"Intercept"`, contains the corresponding bootstrapped intercepts. Each slope and intercept should come from a regression line that predicts column 2 from column 1 of a resample of the given table. The table should have 1 row for each bootstrap replication.

```
[67]: def bootstrap_lines(tbl, num_bootstraps):
      resampled_m = []
      resampled_b = []
      for _ in range(num_bootstraps):
          m, b = fit_line(tbl.sample())
          resampled_m.append(m)
          resampled_b.append(b)
      return Table().with_columns("Slope", resampled_m, "Intercept", resampled_b)

# When you're done, this code should produce the slopes
# and intercepts of 1000 regression lines computed from
# resamples of the faithful table.
regression_lines = bootstrap_lines(faithful, 1000)
regression_lines
```

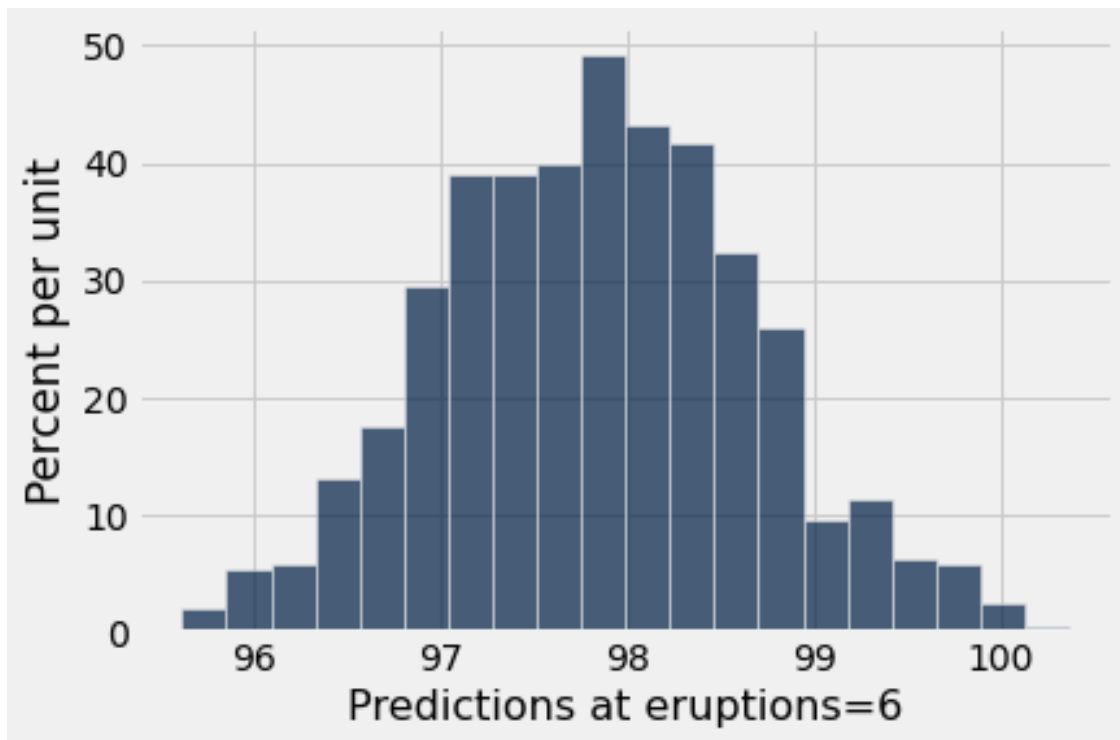
```
[67]: Slope   | Intercept
      10.6263 | 33.916
      10.6813 | 32.9134
      10.8104 | 33.0671
      10.3576 | 34.5081
      10.9507 | 31.9565
```

```
10.4044 | 34.026
10.6981 | 33.1548
10.0846 | 35.9943
10.6442 | 33.602
10.5217 | 33.8853
... (990 rows omitted)
```

Question 3.10 Create an array called `predictions_for_six`. It should contain 1000 numbers. Each number should be the predicted waiting time after an eruption with a duration of 6 minutes, using a different bootstrapped regression line.

Hint: use `regression_lines` from the previous questions.

```
[68]: predictions_for_six = regression_lines["Slope"]*6 +  
      ↪ regression_lines["Intercept"]  
  
# This will make a histogram of your predictions:  
table_of_predictions = Table().with_column('Predictions at eruptions=6',  
      ↪ predictions_for_six)  
table_of_predictions.hist('Predictions at eruptions=6', bins=20)
```



Question 3.11 Create a 95% confidence interval for these predictions.

```
[69]: lower_bound = np.percentile(predictions_for_six, 2.5)
      upper_bound = np.percentile(predictions_for_six, 97.5)

      print('95% Confidence interval for predictions for x=6: (', lower_bound,",",
            ↪upper_bound, '')
```

```
95% Confidence interval for predictions for x=6: ( 96.27441878860185 ,
99.56904740069578 )
```

```
[70]: check('tests/q3_11.py')
```

```
[70]: <gofer.ok.OKTestsResult at 0x7f9351ddd7d0>
```

Question 3.12 Look at the scatter plot of the data at the start of this exercise. Determine which of the following are true, then set `question_12_choice` to an array consisting of the numbers of statements that are true. For example, if you think that 1 and 2 are true but 3 is false, you'd assign `question_12_choice` to be an array consisting of the values 1 and 2.

Statement 1: This confidence covers 95% of waiting times of eruptions in `faithful` that had an eruption duration of 6 minutes.

Statement 2: This interval gives a sense of how much actual wait times differ from your prediction.

Statement 3: The confidence interval quantifies our uncertainty in our estimate of what the true regression line would predict.

```
[71]: question_12_choice = [1, 3]
```

```
[72]: check('tests/q3_12.py')
```

```
[72]: <gofer.ok.OKTestsResult at 0x7f9353931310>
```

1.4 4. Submission

Once you're finished, select "Save and Checkpoint" in the File menu. Your name and course section number should be in the first and last cell of the assignment. Be sure you have run all cells with code and that the output from that is showing.

Double check that you have completed all of the free response questions as the auto-grader does NOT check that and YOU are responsible for knowing those questions are there and completing them as part of the grade for this homework. When ready, click "Print Preview" in the File menu. Print a copy from there in pdf format. (This means you right click and choose print and choose "save as pdf" from your printer options.) You will need to submit the pdf in Canvas by the deadline.

The gopher grader output and/or output from your coding are essential to helping your instructor grade your work correctly and in a timely manner.

Files submitted that are missing the required output will lose some to all points so double check your pdf before submitting.

```
[73]: # For your convenience, you can run this cell to run all the tests at once!
import glob
from gofer.ok import grade_notebook
if not globals().get('__GOFER_GRADER__', False):
    display(grade_notebook('hw11.ipynb', sorted(glob.glob('tests/q*.py'))))
```

Eyeballed error: 13.999999999999998

Aided error: 13.799999999999999

Least-squares regression line: $\text{predicted_y} = -0.6498374231570391 * x + 1.1988619620992735$

95% confidence interval for slope: [10.1053, 11.3227]

95% Confidence interval for predictions for x=6: (96.28274508364186 , 99.50177195394492)

['tests/q1_1.py', 'tests/q1_4.py', 'tests/q1_5.py', 'tests/q1_6.py',
'tests/q2_4.py', 'tests/q2_5.py', 'tests/q2_6.py', 'tests/q2_7.py',
'tests/q2_8.py', 'tests/q3_2.py', 'tests/q3_3.py', 'tests/q3_5.py',
'tests/q3_7.py', 'tests/q3_8.py', 'tests/q3_11.py', 'tests/q3_12.py']

Question 1:

<gofer.ok.OKTestsResult at 0x7f9351dddc0>

Question 2:

<gofer.ok.OKTestsResult at 0x7f9352a0fd0>

Question 3:

<gofer.ok.OKTestsResult at 0x7f935284cb90>

Question 4:

<gofer.ok.OKTestsResult at 0x7f9351cd8450>

Question 5:

<gofer.ok.OKTestsResult at 0x7f9352bf5e10>

Question 6:

<gofer.ok.OKTestsResult at 0x7f9351d22810>

Question 7:

<gofer.ok.OKTestsResult at 0x7f9351b19a90>

Question 8:

<gofer.ok.OKTestsResult at 0x7f9351bc44d0>

Question 9:

<gofer.ok.OKTestsResult at 0x7f9351ba8550>

Question 10:

<gofer.ok.OKTestsResult at 0x7f93527b8a10>

Question 11:

<gofer.ok.OKTestsResult at 0x7f9352173e50>

Question 12:

<gofer.ok.OKTestsResult at 0x7f9352295210>

Question 13:

<gofer.ok.OKTestsResult at 0x7f9352173ed0>

Question 14:

<gofer.ok.OKTestsResult at 0x7f9352295110>


Question 15:

<gofer.ok.OKTestsResult at 0x7f9352297590>

Question 16:

<gofer.ok.OKTestsResult at 0x7f9352282bd0>

1.0



Name: Allan Gongora

Section: 0131